

radare2-1.0

The curse of the milestone

--pancake

Introduction

I'm the author of this thing called `radare` and this talk is gonna be about the state of the project, its weak points, their fixes and the future plans.

DevRevs: A concept that I invented to refer to the reverse engineers that are also developers (and vice versa). Main contributors have this profile.

If you enjoy this talk you are probably one of those.

The culture of versioning

The version number of a program can tell you a lot of things about the software and the developers behind. There are several ways of understanding version numbers, and this may vary depending on the project or the language used.

What do 1.0 means?

- Initial release (others use 0.0.1 for this)
- First stable version
- Complete and finished version of the software



The culture of versioning

Version numbers can be just a number

- 43 (browsers like Mozilla Firefox or Google Chrome)
- 1.x (nmap, coreutils, jq, ...)
- 1.x.x (like r2, capstone, git, gnome, apache, ...)
 - Maybe it encodes abi compatibility level.
- 1.x.x.x (tor, sqlite, snort, ...)
- 1a (like openssl.)
- 1-beta
- YYYYMMdd
- YYYY-MM
- Snapshot
- git
- ...



ABI compatibility

According to the semantic versioning: (1)

- MAJOR version when you make incompatible API changes.
- MINOR version when you add functionality in a backwards-compatible manner.
- PATCH version when you make backwards-compatible bug fixes.



*1 <http://semver.>

Timeline

Some understand versions as an evolution line of the project, where major number is used to represent the milestones that show big changes in the interface or the program functionalities.

Sometimes two big version numbers are developed in parallel. Version numbers increase every some specified amount of time.

GTK changed their versioning and release cycles recently, most projects are moving to more frequent releases.


6 week release cycle since 0.10 (like Rust). Frequent updates, aims stable git and fixed milestones handy for project management.

Branches

Some projects use different branches for devel, which requires an extra effort to backport patches to every branch in order to construct proper LTS updates.

In r2land we use 1 branch:

- Evolve faster
- Always refactoring
- Less maintenance work
- Focus on features and stability
- Drop deprecated APIs and commands ASAP
- Using pull requests with jenkins+travis+appveyor integration



Are we 1.0 yet?

Many people thought that it would arrive after 0.9...

But 0.9.1 arrived

After 0.9.9 many people believed the next would be finally 1.0...

But 0.10 appeared

Last release is 0.10.5, guess what will be next...

r2-v1.0

Yes, 1.0 has arrived, and this doesn't mean that it's a stable or complete and finished piece of software.

What 1.0 means for radare2 is that it has reached 10 years!



r2-v1.0

The first release of 2017 will be the r2-1.0

- += 0.1 every 6 weeks
- += 1.0 every year

This way the versioning represents different snapshots over time, rather than informing false abi compatibility information.

Development Model

In r2land we used to not follow any strict development model, which turned into a continuous evolution of the practices that worked for us, development, the packagers and the users.

- Using ACR (configure+make)
- sys/install.sh to simplify updates
- Handmade makefiles
- Precompiled sdb databases.
- ...

I used to increase a 0.0.1 increment with about 1000 commits after some months, so version.

Weak points

...



Weak points = Complains

Those are the main complains from the users:

- Commands are complex and hard to remember
- Nobody knows all features of r2
- It have no GUI
- Doesn't compiles
- Bindings are not idiomatic and hard to install/use
- There's no documentation
- It's slow when working with big files
- Not usable for large projects
- It is vulnerable because it's in C



Problem: Huge Codebase

It is true that r2 is quite big in size, this can be a problem for embedding into new targets with low resources, or when compiling multiple versions to find which commit broke a regression test.

Solutions

- Always refactor, removing dead or wrong code, optimizing
- Disable plugins with `./configure-plugins (sys/tiny.sh)`
- Move code from r2 to r2-extras and make it available via r2pm
- Specify default plugins
- Use CI systems to delegate
- Do more review on new code
- Standardize syntax and indentation rules

Dependencies

Aim to be dep-free, but that's because it ships them inside, and some are optionally compilable against system libraries:

- **Binutils** stuff (GPL) and not available as a library
- **Capstone** (last release is old and it comes with some vulns)
- **LibMagic** (GNU one is vulnerable, so I took OpenBSD one (which source is crap, but at least smaller, and fixed a couple of bugs (reported upstream))
- **TCC** - Tiny C Compiler (stripped and modified to fit into r2)
- **GRUB** (GPL) - FileSystem and Partitions (full of vulns, fixed many of them + api)
- **SDB** - Done by me, and kept in sync into the r2 codebase
- **Udis86** - handy for small builds, but abandoned (capstone++)
- **(Libre|Open|Boring)SSL** - Used for BigNum+SSL (unused)

Vulnerabilities and Patches

Sometimes external dependencies have vulnerabilities, found by us or just reported somewhere else. In r2land, we handle segfaults and vulns as a maximum priority issue. Which use to be fixed before a day of being noticed.

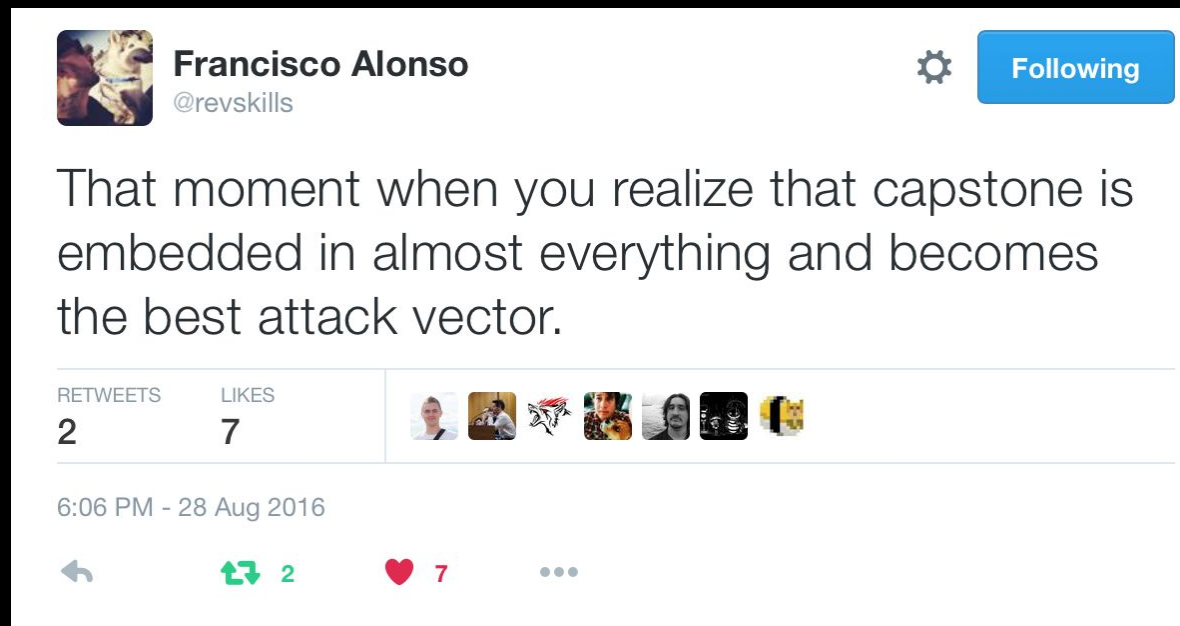
We ship also some patches to fix missbehaviours.

- aim to get pushed upstream
- use embedded libs + patches by default
- pick a specific branch + commit to detect regressions

Vulnerabilities In Capstone

Revskills CVEs affecting git version of Capstone.

- CVE-2016-7151 <https://github.com/aquynh/capstone/pull/725>
- CVE-2016-4044 <https://github.com/aquynh/capstone/issues/729>
- CVE-2016-3160 <https://github.com/aquynh/capstone/issues/730>



Dependency Problems

- We need to keep in sync with master
 - Not always easy if patches are applied
 - Patches on top of those projects must be simple
- We must push all fixes upstream, as well as report issues
- Distros dislike static linking, so we need to provide a way to link against system ones
 - Can make r2 vulnerable, or behave different than using the shipped versions
 - Simpler bug fixing for distros
 - Smaller Builds

Solution: Packaging

Since r2-0.10.0 (2015-09-24) there's a program called r2pm that allows to install and manage external dependencies and plugins, installing everything in the home or system (-g)

(DEMO)



C Language

C is the most portable language out there, there are compilers that can generate code for microcontrollers, or even to javascript to rule the web.

The good side

- It is pretty low level. Mostly WYWIWYG
- Many analysis and helper tools available
- It's easy to debug

The bad side

- Requires skilled developers
- More time invested in code reviews
- Some time of the dev must be reinvested in fixing bugs

Solution: Code Review

C is a weak language in the sense that provides little abstraction from the generated code and requires the developer to handle memory, resources, buffer sizes, endians swaps and other caveats that can make spend more time fixing bugs than developing.

- Multiple reviews before merging helps
- Defining good code practices
- Using Scan Coverity and Clang Analyzer
- Fuzzing and



Coding Style

We use a custom coding style inspired by GNOME and Perl standards.

No tool is able to indent C code properly.

- Avoids control flow bugs (goto fail)
- Makes the code more readable
- Slowly adopting the coding style
- Not reindenting 3rd party code

RDD - Regression Driven Development

Some projects use TDD, which means that you write tests before code.

In r2land I invented the RDD, which is basically the same, but in reversed order. The reason for this is because the way r2 is developed and its culture.

- Most parts of r2 has been done for a specific need or fun.
- radare->radare2 transition showed that having a perfect design from the beginning is impossible for projects big like this (and even small like dwm)
- Moving fast is a key piece of the project, we want things now.
- Writing tests is boring
- Testing everything is impossible (code coverage)
- Integration with Travis is important to avoid regressions, slowdowns the devel (wait for results), but avoids mayhem.

RDD - Regression Driven Development


DEMO: Write a test

```
$ git clone https://github.com/radare/radare2-regressions
```

```
$ cd radare2-regressions
```

```
$ vim t/my_test
```

```
$ sh run_tests.sh t/my_test
```



Infrastructure

Having a testsuite is not enough for catching the bugs in time, it is very useful to have travis building every single commit and running the testsuite on different system setups.

- Testsuite only running on Linux-x86-64 and OSX-x86-64 (travis+jenkins)
- Automated Windows, Android builds
- Manual iOS and Android updates
- Releasing requires some time (6 week helps)

Endianness

R2 is the process of being fully big-endian friendly. (thx damo22)

The host endianness is now ignored completely, and just uses binary operations to grab the values in the specified endian (big or little).

Configurable via `cfg.bigendian`.

- Host Endian (local cpu)
- Target Endian (target `asm.arch` for `asm/disasm` and `pxw/pointers`)
- Implicit Endian (bin files with fixed endianness)

Memory leaks

Currently it is a problem. But not really dramatic.

- Use valgrind to profile the heap usage
- Commands with loops shouldn't leak (pd, ...)
- To leak or to double free
- Ownership concept problems
- No reference counting



Other implicit bugs from C

- Null dereferences
- Buffer Overflows
- Heap overflows / underflows
- Format strings
- Use After Free
- Pointer alignment issues
- Double Frees
- Corrupted pointers
- Memory leaks
- OOB reads
- ...



Alternative languages

The only language that fit in the field of C is probably Rust.

- No big or mandatory runtime library
- No garbage collector
- Low level
- Easily to talk

Bad points:

- Young, but stable.
- Hard to find experienced programmers
- Mind shift.



Other languages

As long as r2 is modular and pluggable it is possible to rewrite every single module in a different language being able to reuse the other parts and plugins written in C and just replacing the specific library.

- Swift
- Vala
- Nim/Ooc

Extending or replacing parts of r2 in dynamic or scripting languages like Python, NodeJS, Ruby or C# is possible, but for Core, it's preferably to choose a language that:

- No garbage collection
- Low level (no mangling for public symbols, no runtime)
- On top of POSIX (slightly blocking for kernel or embedded land)

Performance

Many people think that r2 must be fast because it is written in C. This is not a true statement because the way memory is managed and the algorithms implemented it can result in slower than other languages. But anyway, r2 is reasonably fast for most use cases, except for:

- Big files
 - Having a large number of flags or functions results in heavily slow r2
 - IO is complex
 - Many memory copies
 - RBin parsing duplicates memory
 - Repeated processes (analysis in pd, etc)
- Remote debugger
 - Non-cached requests

Performance: IO

The `libr/io` is responsible of abstracting the access to files, memory, processes. It exposes several layers to define sections, maps, specify permissions, caching temporal patches, non-memoized functions, etc. All this complexity results in slow read operations.

- Do single reads in a block before a processing loop
- Don't do small reads on every iteration
- Having many sections

SDB

The disk/memory string key=value database is used in many places of r2 for:

- Easy to use hashtables
- Can be mmaped from disk, so it takes no memory
- Inspectable by the user (k command)
- Leak-free and clean api

The bad side

- A hashtable is not always
- Strings are not always the more optimal way to represent stuff
- Used for storing data of projects

Analysis

Statically analyzing code is not an easy problem to solve.


- Same loop for all archs
- Each anal plugin provides basic backend logic
- Merge with anal + asm?
- Reuse information from disasm and analysis to avoid having to disasm every instruction twice
- Recursive analysis is weak
 - Break analysis with wrong code
 - Code coverage is hard
- Requires emulation sometimes
- Storage problem



Parallelism/Threading

R2 is monotask, it uses no threads at all. This decision is on purpose.

- Simple design
- Avoid unnecessary mutexes
 - Probably deadlocks
- Co-routines and async are hard
- Still lot to optimize before going into this
- Queues and snapshotting sounds better to me
- Basic support for threads (background webserver, ..)



Debuggers

R2 is monotask, it uses no threads at all. This decision is on purpose

Weak points

- Unstable and poorly tested
- Early threading support
- All are sync operations
- More support for kernel
- Port dmh to more platforms



Bindings

Bindings are automatically generated by valabind. A tool I wrote that reads vapi descriptions and generates bindings for swig or other targets.

Those bindings are not idiomatic and are pretty much 1:1 with the C API, so the same API documentation can be reused across all the languages.

Imposes some restrictions in the C code base of r2 to integrate with Vala:

- Don't expose unions
- RList everywhere
- Implicit ownership rules
- Naming conventions for functions to be OO-like

r2pipe

Simplest interface possible with r2: the shell. It provides a way to run a command and get the output in return.

This simple concept can be used thru sockets, file descriptors, native apis via `dlopen()`, http, rap, etc..

Deserializing JSON objects is faster than FFI.

Easy to port to new languages.


```
// demo
```



User Interface

The main complain from new users is the lack of a user interface.

- Writing a UI is boring
- R2 api evolves and requires constant updates in both sides
- The shell is more powerful and expressive
- Need ways for remoting



Solution 1: TUI

```
$ npm install blessr2
```

```
blessr2 /bin/ls @ entry0
0x100001964 00 0099 .....
0x100001a64 01 0085 .....
0x100001b64 02 00b3 .....
0x100001c64 03 00b3 .....
0x100001d64 04 00b5 .....
0x100001e64 05 00b9 .....
0x100001f64 06 00ad .....
0x100002064 07 0089 .....
0x100002164 08 00ad .....
0x100002264 09 00a6 .....
0x100002364 0a 00a7 .....
0x100002464 0b 00a7 .....
0x100002564 0c 00af .....
0x100002664 0d 00aa .....
0x100002764 0e 00ae .....
0x100002864 0f 00ae .....
0x100002964 10 00ae .....
0x100002a64 11 00b3 .....
0x100002b64 12 00b9 .....

rax 0x00000000 rdx 0x00000000 r8 0x00000000 r11 0x00000000 r14 0x00000000 rbp 0x00000000
rbx 0x00000000 rs1 0x00000000 r9 0x00000000 r12 0x00000000 r15 0x00000000 rsp 0x00000000
rcx 0x00000000 rdi 0x00000000 r10 0x00000000 r13 0x00000000 rip 0x00000000 rflags

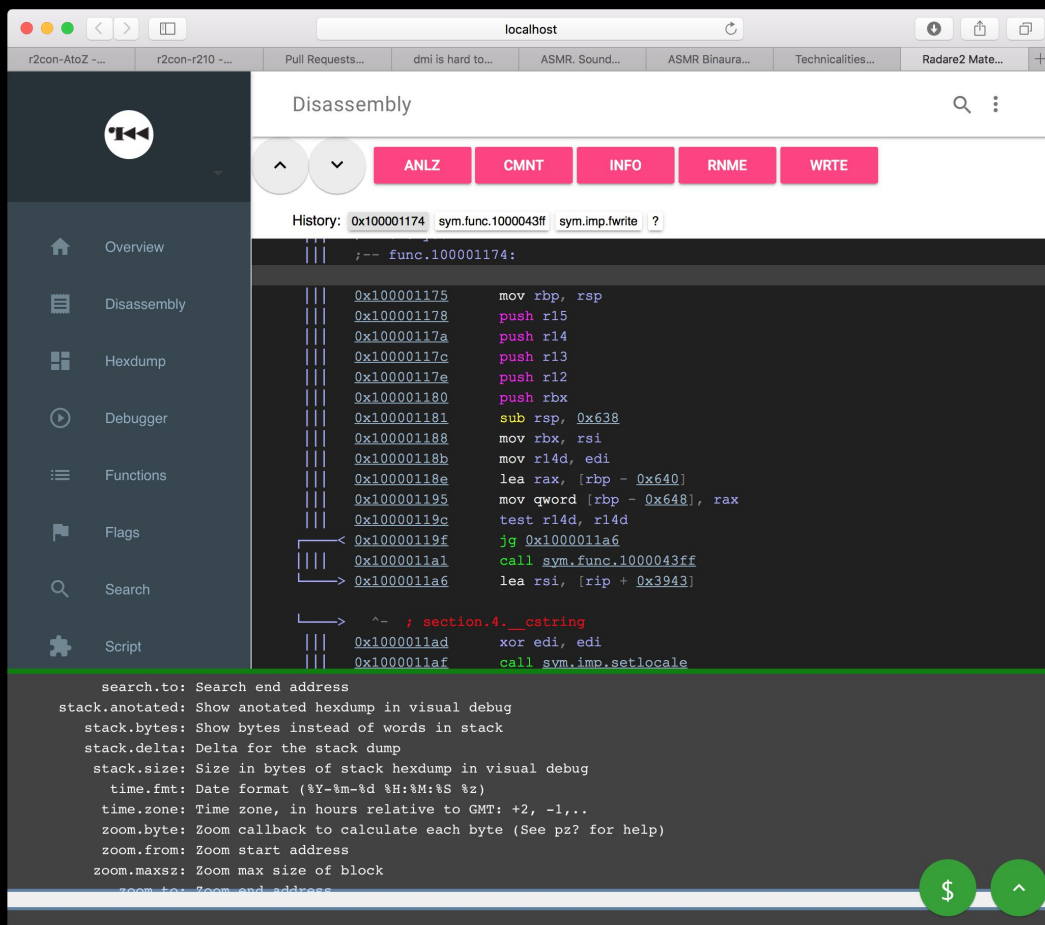
0x100003564 1c 009c .....
0x100003664 1d 00b5 .....
0x100003764 1e 00ad .....
0x100003864 1f 00b1 .....
0x100003964 20 00b5 .....
0x100003a64 21 00b9 .....
0x100003b64 22 00b2 .....
0x100003c64 23 00aa .....
0x100003d64 24 00a9 .....
0x100003e64 25 00a9 .....
0x100003f64 26 00ae .....
0x100004064 27 00a5 .....
0x100004164 28 00a7 .....
0x100004264 29 00ab .....
0x100004364 2a 00a5 .....
0x100004464 2b 00a5 .....
0x100004564 2c 0072 .....
0x100004664 2d 0077 .....
0x100004764 2e 0081 .....
0x100004864 2f 009e .....
0x100004964 30 00a0 .....
0x100004a64 31 00b3 .....
0x100004b64 32 0098 .....
0x100004c64 33 0097 .....
0x100004d64 34 008d .....
0x100004e64 35 009c .....
0x100004f64 36 0058 .....
0x100005064 37 0049 .....
0x100005164 38 0049 .....

0x100001c64 0000 488d 3db4 2e00 0048 8d35 b42e 0000 ..H.....H.5....
0x100001c74 e83f 2800 0084 c00f 8425 0100 00e9 1901 .?(.....%.
0x100001c84 0000 4183 fe02 7c19 498b 7530 31c0 488d ..A...|.1.u01.H.
0x100001c94 3d03 2f00 00e8 ce28 0000 c605 7438 0000 =/.....(.....t8..
0x100001ca4 014c 89e7 488b 75c8 e831 2800 0049 89c6 .L..H.u..1(.I..
0x100001cb4 488d 3d66 2e00 0048 8d35 662e 0000 e8f1 H.=f...t&4.u"M..L
0x100001cc4 2700 0045 85ff 7426 3401 7522 4d85 f64c '.E...t&4.u"M..L
0x100001cd4 89f0 741a 0fb7 4858 83f9 0d75 0840 c740 ..t...HX...u.H.@
0x100001ce4 1801 0000 0048 8b40 1048 85c0 75e6 4c89 .....H.@..H..u.L
0x100001cf4 ef4c 89f6 e88f 0100 004d 85f6 448b 75d4 .L.....M...D.u.
0x100001d04 0f84 9c00 0000 8a05 f137 0000 3401 a801 .....7...4....
0x100001d14 0f84 8c00 0000 ba04 0000 004c 89e7 4c89 .....L...L.L.
0x100001d24 eee8 d027 0000 eb7a 410f b745 5883 f80c .....zA..EX...
0x100001d34 7f1a f1c8 83f8 0677 6948 6904 8348 01d8 .....wHc...H..
0x100001d44 ffe0 418b 7d38 4983 c568 eb33 83f8 0d75 ..A.)81...h.3...u
0x100001d54 5148 8d3d c52d 0000 488d 35c5 2d00 00e8 QH.=...H.5....
0x100001d64 5027 0000 4585 ff74 3934 0175 3541 8b7d P'..E..t94.u5A.}
0x100001d74 3849 83c5 6885 ff7b 0200 0000 0f44 f8e8 8I..h.....D..
0x100001d84 2c28 0000 4889 c131 c048 8d3d fb2d 0000 ..(H...1.H.=...-
0x100001d94 4c89 ee48 89ca e845 2800 00c6 0572 3700 .L..H...E(.....r7.
0x100001da4 0001 4c89 e7e8 4627 0000 4989 c54d 85ed ..L...F'..I..M..

-- main:
-- entry0:
-- func.100001174:
0x100001174 push rbp
0x100001175 mov rbp, rsp
0x100001178 push r15
0x10000117a push r14
0x10000117c push r13
0x10000117e push r12
0x100001180 push rbx
0x100001181 sub rsp, 0x638
0x100001188 mov rbx, rs1
0x10000118b mov r14d, edi
0x10000118e lea rax, [rbp - 0x640]
0x100001195 mov qword [rbp - 0x640], rax
0x10000119c test r14d, r14d
0x10000119f jg 0x1000011a6
0x1000011a1 call sym.func.1000043ff
0x1000011a6 lea rs1, [rip + 0x3943] ; 0x100004af0 ; section.4.__cst
0x1000011ad xor edi, edi
0x1000011af call sym.imp.setlocale
0x1000011b4 mov r12d, 1
0x1000011b8 mov edi, 1
0x1000011bf call sym.imp.isatty
0x1000011c4 test eax, eax
0x1000011c6 je 0x100001229
0x1000011c8 mov dword [rip + 0x42fe], 0x50 ; [0x1000054d0:4]=00 LEA
0x1000011d2 lea rdi, [rip + 0x3918] ; 0x100004af1 ; str.COLUMNS ;
0x1000011d9 call sym.imp.getenv
0x1000011de test rax, rax
```


Solution 2: WebUI

```
$ r2 -c=H /bin/ls
```



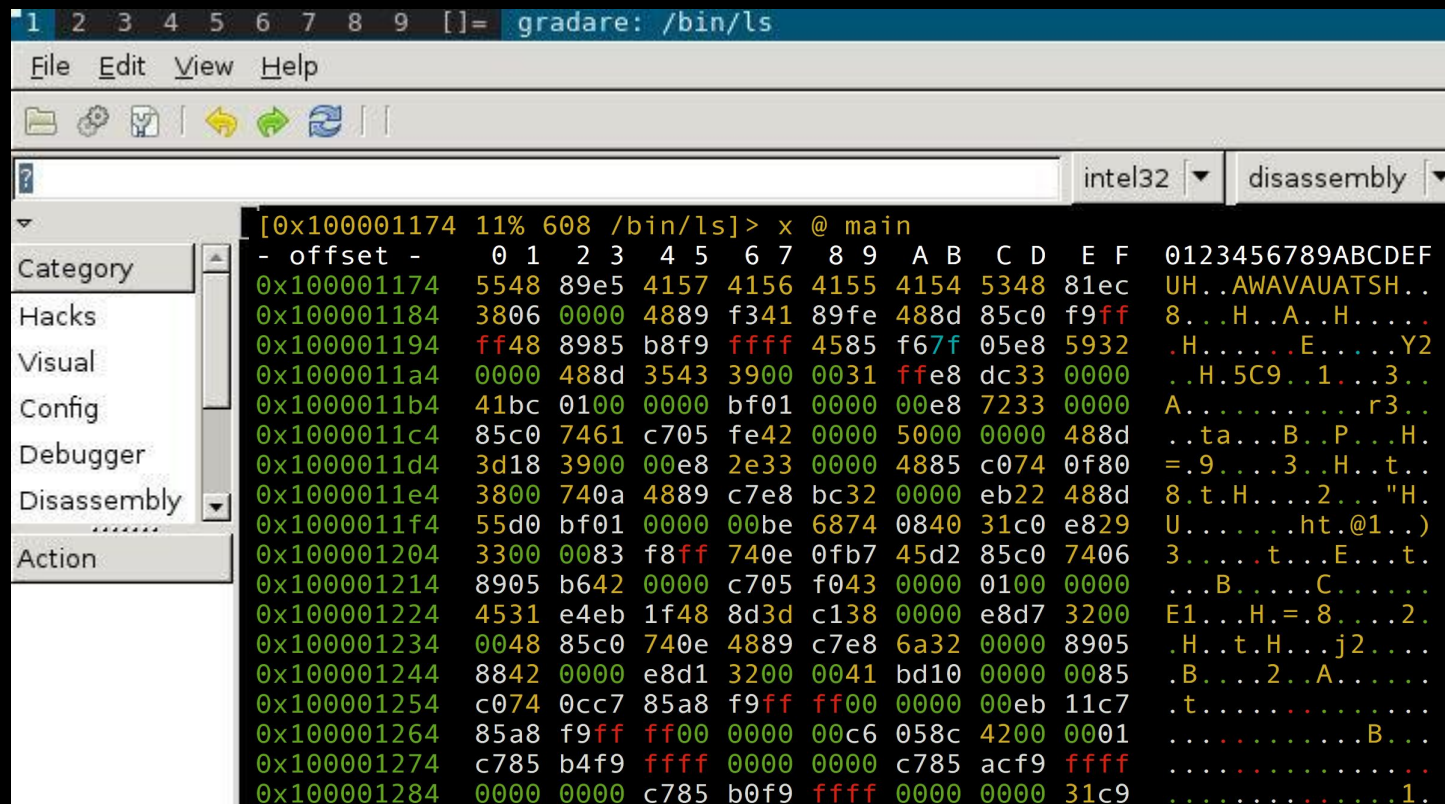
Solution 3: Native App

I have running versions of r2 on iWatch, unjailbroken iPhones, and wrote a couple of GUIs in GTK and QT. Hey! But there's also Bokken and Gradare2!

- Boring to develop, lack of widgets and task queue
- Jumping up and down from core to gui is tiring
- Better user experience

Solution 3: Native App: Gradare2

Rewrite of gradare, wip gtk3, vte-based UI

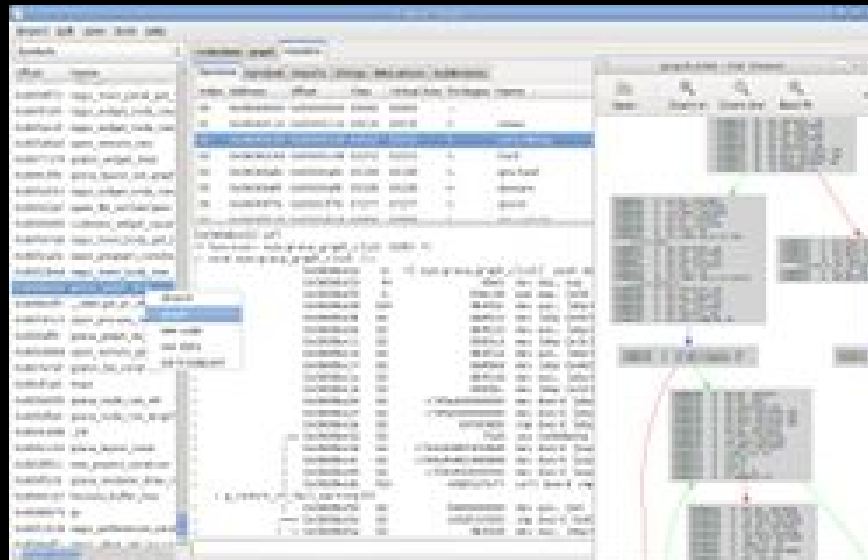


The screenshot shows the Gradare2 application interface. At the top, a command line displays 'gradare: /bin/ls'. Below this is a menu bar with 'File', 'Edit', 'View', and 'Help'. A toolbar with various icons is positioned below the menu bar. The main window is divided into two panes. The left pane contains a sidebar with a 'Category' list (Hacks, Visual, Config, Debugger, Disassembly, Action) and a search bar. The right pane displays a disassembly view for the 'intel32' architecture. The disassembly shows memory addresses, offsets, and hex values, along with a corresponding ASCII representation of the output. The output appears to be a list of files and directories, including 'UH..AWAVAUATSH..', '8...H..A..H....', '.H.....E.....Y2', '..H.5C9..1...3..', 'A.....r3..', '..ta...B..P...H.', '=..9....3..H..t..', '8..t.H....2...\"H.', 'U.....ht.@1..)', '3.....t...E...t.', '...B.....C.....', 'E1...H.=.8....2.', '.H..t.H...j2....', '.B....2..A.....', '.t.....', '.....B...', '.....', and '.....1..'. The disassembly view is titled 'intel32' and 'disassembly'.

```
1 2 3 4 5 6 7 8 9 []= gradare: /bin/ls
File Edit View Help
[Icons]
[?] intel32 disassembly
Category
Hacks
Visual
Config
Debugger
Disassembly
Action
[0x100001174 11% 608 /bin/ls]> x @ main
- offset - 0 1 2 3 4 5 6 7 8 9 A B C D E F 0123456789ABCDEF
0x100001174 5548 89e5 4157 4156 4155 4154 5348 81ec UH..AWAVAUATSH..
0x100001184 3806 0000 4889 f341 89fe 488d 85c0 f9ff 8...H..A..H....
0x100001194 ff48 8985 b8f9 ffff 4585 f67f 05e8 5932 .H.....E.....Y2
0x1000011a4 0000 488d 3543 3900 0031 ffe8 dc33 0000 ..H.5C9..1...3..
0x1000011b4 41bc 0100 0000 bf01 0000 00e8 7233 0000 A.....r3..
0x1000011c4 85c0 7461 c705 fe42 0000 5000 0000 488d ..ta...B..P...H.
0x1000011d4 3d18 3900 00e8 2e33 0000 4885 c074 0f80 =..9....3..H..t..
0x1000011e4 3800 740a 4889 c7e8 bc32 0000 eb22 488d 8..t.H....2...\"H.
0x1000011f4 55d0 bf01 0000 00be 6874 0840 31c0 e829 U.....ht.@1..)
0x100001204 3300 0083 f8ff 740e 0fb7 45d2 85c0 7406 3.....t...E...t.
0x100001214 8905 b642 0000 c705 f043 0000 0100 0000 ...B.....C.....
0x100001224 4531 e4eb 1f48 8d3d c138 0000 e8d7 3200 E1...H.=.8....2.
0x100001234 0048 85c0 740e 4889 c7e8 6a32 0000 8905 .H..t.H...j2....
0x100001244 8842 0000 e8d1 3200 0041 bd10 0000 0085 .B....2..A.....
0x100001254 c074 0cc7 85a8 f9ff ff00 0000 00eb 11c7 .t.....
0x100001264 85a8 f9ff ff00 0000 00c6 058c 4200 0001 .....B...
0x100001274 c785 b4f9 ffff 0000 0000 c785 acf9 ffff .....
0x100001284 0000 0000 c785 b0f9 ffff 0000 0000 31c9 .....1..
```

Solution 3: Native App: Ragui

Ragui - never published full native GUI written in Vala and GtkON 6 years ago.. In the works to be ready for full GTK3



Solution 3: Native App

\$ r2pro

The screenshot displays the r2pro application interface, which is a front-end for Radare2. The main window is titled "main" and shows the assembly code of a function. The code is organized into sections, with the current section being "section_end.plt:". The assembly code is shown in a list view on the left, with the corresponding disassembly on the right. The disassembly includes comments and cross-references to other parts of the program.

On the right side of the interface, there is a control flow graph (CFG) showing the flow of execution between different blocks of code. The graph is a directed acyclic graph (DAG) with nodes representing basic blocks and edges representing control flow. The nodes are labeled with addresses and instructions.

Below the CFG, there is a "Function: main" panel showing the function's information, including its name, arguments, and return type. The panel also displays the function's offset information, such as the family, stack, and size.

At the bottom of the interface, there is a "Sections" panel showing the list of sections in the program, including their names, sizes, and addresses. The sections are listed in a table with columns for Name, Size, Address, and End Address.

The interface also includes a "Dashboard" panel at the bottom left, which provides a summary of the program's analysis, including the number of instructions, the number of basic blocks, and the number of functions.

Questions?
